



What the joint cohort model actually does

It estimates an aggregate booking-to-revenue bridge. It does not observe individual reservation histories.



Same question, two forecasting rules




“Joint” means one consistent allocation supports both views: where this quarter’s revenue came from, and where a booking quarter’s allocated fees go. Many allocations can fit the same total revenue.

Weights shown are the current fitted predictor weights, not observed booking shares or conversion probabilities.
Seasonal multipliers combine monetization, timing, cancellations, mix and specification error; RNPL is not separately identified.


Accuracy depends on the target and the sample

Raw midpoint RMSE; lower is better. Reconstructed historical forecasts, not live performance.


Guide announcement 1 ahead | n=10




Guide announcement 2 ahead | n=10



Guide announcement 3 ahead | n=9




The favorable p+4 result is sample-sensitive



Top/left panels: W2 common rows, targets from 2024 onward. Each horizon has its own eligible sample.
Bottom right: fixed lag versus guide growth at p+4. Its nine-row win fails when earlier eligible quarters are restored.
All six model-horizon promotion claims fail the broader candidate-eligibility audit; failure is not proof of uselessness.

A revenue advantage does not yet establish a guidance trade

Same calendar origin, same ten W2 quarters, same target within each panel. Street = the held DoltHub series.




Live Q4 illustration: joint guide \$3.185bn versus revenue consensus \$3.200bn looks negative.
Against the same-cushion implied guide (\$3.142bn), it is +\$43.6m. Those comparisons answer different questions.

W1 joint revenue RMSE is also lower: \$57.4m versus \$89.9m (n=11). The favorable finding is retained.
Historical Street snapshots have inherited provenance; the 14 used vintages remain uncertified at row level.
The Street guide proxy is revenue consensus divided by our cushion. It is not an observed expectation of management’s guide.

Three future guides: a short trade can still rely on forecast inputs

As of 15 September 2026; latest company actuals published 6 August. Q3 guidance is already issued.




Booking carry is economically real. Our reported aggregate data do not measure its full size.
In the live fitted rules, Q1/Q2 estimates depend entirely on projected GBV inputs.

Points are conditional research forecasts, not promoted calls. Lines between quarters do not remove seasonality.
Known-GBV share measures exposure of the model’s predicted dollars to published GBV; it is not actual revenue already booked.
There are no calibrated next-event prediction intervals, and no quarterly Street series here for Q1/Q2 2027.


The most useful repair is better future-GBV information

An unavailable perfect-input diagnostic identifies room to improve; the tested flight repair does not deliver it.

Perfect GBV inputs reduce fixed-model error



The specific flight remedy fails its gate



Prioritize booking/stay timing, cancellation outcomes and correctly dated quarterly signals.
More fitted cohort weights cannot manufacture the missing information.

Left: same W2 common samples within each horizon (n=10,10,9). Oracle uses actual future GBV and is impossible at trade time.
Right: the preregistered calendar-flight test uses nine identical W1/W2 rows and one fitted slope; only two current-quarter features.
Error components covary; the oracle improvement is conditional potential, not an achievable guarantee or a causal decomposition.